

The Window Valet

Step-by-step breakdown for the final two phases of the original roadmap

Prepared for Josh LeClair by John Carlo Caintic · Marketing Strategy & Automation · August 29, 2026

Weeks 1 through 6 are either done or already covered under the current Week 3-4 roadmap (scheduling research and the Mapsly prototype folded into that list). This covers what's left: the Partner system and AI/advanced automation, the two phases of the original 8-week plan that haven't started yet.

7. Partner System

Where this stands today: the "3. Partner Pipeline" already exists and is clean (New Partner Lead → Contacted → Negotiating → Finalized Partnership → Closed-Lost, no duplicates found in the audit). A "Partner New Leads Opt-in" workflow is already live and enrolling partners. What's missing is a dedicated nurture sequence built specifically for partners, separate from the end-customer Drip/Nurture engine, since a referral partner relationship needs different content and pacing than a homeowner deciding on window treatments.

Blocker: this can't start until tag standardization gets Josh's sign-off. The current "warm"/"warm lead" tag naming is inconsistent, and the nurture branch logic below depends on checking that tag reliably, which isn't possible until it's collapsed into one consistent tag across the account.

- 1. Resolve tag standardization first.** Collapse the "warm"/"warm lead" naming cluster into a single, consistent tag, and get Josh's sign-off on the final taxonomy before anything else here starts.
- 2. Audit "Partner New Leads Opt-in."** Understand exactly what this existing workflow currently does, it's applying an interim "Drip" tag as a stopgap right now, before changing anything.
- 3. Define what makes partner content different.** Partners are referral relationships, not homeowners deciding on a purchase. Messaging should focus on the partnership itself and referral incentives, not consultation booking language.
- 4. Design the branch logic.** Once the tag is standardized: check whether a partner is tagged "Warm Lead," branch to a faster, more personal cadence for warm partners versus a slower, standard cadence for cold ones.
- 5. Tie it to the existing Partner Pipeline stages.** The nurture sequence should exit cleanly once a partner progresses (Contacted, Negotiating, Finalized Partnership), same exit-on-progress pattern already proven in the end-customer Drip/Nurture workflow.
- 6. Build it reusing the proven mechanics.** Same wait-step and exit-check pattern from Drip/Nurture, since that structure already works, this isn't a case for reinventing the approach.
- 7. Test with a real or sample partner contact end-to-end** before turning it on for the full list.

8. **Retire the interim "Drip" tag stopgap** on "Partner New Leads Opt-in" once the real nurture sequence is live and confirmed working.
-

8. AI & Advanced Automation

Where this stands today: nothing scoped yet, this was always meant to be the final phase once everything else was stable, which it now largely is. This phase is discovery-first, not build-first, the goal below is to find and prioritize real opportunities, not commit to a specific build before knowing what's worth automating.

1. **Inventory what AI capability already exists in the account.** GHL includes native AI Studio, AI Agents, and Conversations AI Agent features, confirm what's active, what's unused, and what's actually available on the current plan.
2. **Identify repetitive, manual work as candidates.** Look specifically at places a human is doing something predictable and time-consuming: drafting first responses, qualifying leads, personalizing proposal follow-ups.
3. **Decide on first-response AI Agents specifically.** We already saw the SLA settings include a toggle for how AI Agent replies count toward response time, worth deciding intentionally whether to pilot AI-handled first responses rather than leaving that setting as an afterthought.
4. **Prioritize by effort versus impact.** Quick wins first (AI-assisted message drafting, proposal follow-up personalization) before bigger lifts (a full AI qualification or booking agent).
5. **Pick one pilot, define success criteria up front.** What does "working" actually look like for that one use case, response time, conversion rate, hours saved, before testing it live.
6. **Run the pilot, then decide.** Document what worked and what didn't, and use that to decide what expands account-wide versus what gets dropped.

Part of the ongoing GHL stabilization work under the original 8-week roadmap.